

Hyperparameters for `LogisticRegression` (Scikit-Learn)

Logistic Regression has several hyperparameters that control its behavior, regularization, and optimization. Below is a **comprehensive list** of key hyperparameters, their purposes, and typical values.

1. Core Hyperparameters

Hyperparameter	Description	Common Values
<code>penalty</code>	Type of regularization (<code>'l1'</code> , <code>'l2'</code> , <code>'elasticnet'</code> , <code>'none'</code>).	<code>'l2'</code> (default)
<code>C</code>	Inverse of regularization strength (smaller = stronger regularization).	<code>1.0</code> (default), <code>0.01</code> – <code>10</code>
<code>solver</code>	Optimization algorithm (see solver compatibility).	<code>'lbfgs'</code> (default)
<code>max_iter</code>	Maximum number of iterations for solver convergence.	<code>100</code> (default), <code>100</code> – <code>1000</code>
<code>class_weight</code>	Weights for classes (handles imbalance). Use <code>'balanced'</code> or custom dict.	<code>None</code> (default), <code>'balanced'</code>

2. Regularization-Specific

Hyperparameter	Description	Notes
<code>l1_ratio</code>	Mixing parameter for <code>'elasticnet'</code> penalty (0 = L2, 1 = L1).	Only if <code>penalty='elasticnet'</code>
<code>dual</code>	Dual or primal formulation (for <code>penalty='l2'</code> and <code>solver='liblinear'</code>).	Rarely used (<code>False</code> default)

3. Solver-Specific Options

Solver (<code>solver</code>)	Supported Penalties	Use Cases
<code>'lbfgs'</code> (default)	<code>'l2'</code> , <code>'none'</code>	Small-to-medium datasets, multiclass.
<code>'liblinear'</code>	<code>'l1'</code> , <code>'l2'</code>	Small datasets, binary classification.

Solver (<code>solver</code>)	Supported Penalties	Use Cases
'newton-cg'	'l2', 'none'	Medium datasets, multiclass.
'sag' / 'saga'	'l1', 'l2', 'elasticnet'	Large datasets ('saga' supports L1).

4. Multiclass Settings

Hyperparameter	Description	Options
<code>multi_class</code>	Strategy for multiclass problems ('auto', 'ovr', 'multinomial').	'auto' (default)
<code>intercept_scaling</code>	Scales the intercept (useful for <code>solver='liblinear'</code> with <code>penalty='l2'</code>).	1.0 (default)

5. Advanced Tuning

Hyperparameter	Description	Example Values
<code>tol</code>	Tolerance for stopping criteria (optimization precision).	1e-4 (default), 1e-3
<code>fit_intercept</code>	Whether to fit an intercept term.	True (default)
<code>verbose</code>	Log output during training (useful for debugging).	0 (silent), 1 (verbose)
<code>warm_start</code>	Reuse previous solution for incremental training.	False (default)

Example: Hyperparameter Tuning with GridSearchCV

pythonCopy

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import load_iris

# Load data
X, y = load_iris(return_X_y=True)

# Define hyperparameter grid
param_grid = {
    'penalty': ['l1', 'l2'],
    'C': [0.01, 0.1, 1.0, 10],
    'solver': ['liblinear', 'saga'],
```

```
'max_iter': [100, 200]
}

# Search best hyperparameters
model = LogisticRegression()
grid_search = GridSearchCV(model, param_grid, cv=5)
grid_search.fit(X, y)

print("Best hyperparameters:", grid_search.best_params_)
```

Key Takeaways

1. Regularization:

- Use `penalty='l1'` for feature selection (sparse solutions).
- Use `penalty='l2'` for general-purpose regularization.

2. Solver Choice:

- `'liblinear'` for small datasets, `'saga'` for large datasets or L1.

3. Class Imbalance:

- Set `class_weight='balanced'` if classes are uneven.

4. Multiclass:

- Prefer `multi_class='multinomial'` for better multiclass accuracy.

For production, log hyperparameters in MLflow:

```
python

import mlflow

with mlflow.start_run():
    mlflow.log_params(grid_search.best_params_)
    mlflow.sklearn.log_model(grid_search.best_estimator_, "model")
```

Copy